

Biomedical Data Science & AI

Exercise sheet 2 - Introduction - Due date: May 4th

Submitted to:

Mohamed Aborageh : s0moabor@uni-bonn.de

Vinay Srinivas Bharadhwaj: s0vibhar@uni-bonn.de

Yasamin Salimi: yasisali@uni-bonn.de

Exercise 1 - Understanding Your Dataset (total: 13 points)

Load the processedClevelandData.csv dataset. The features for the dataset are described in the featureDescription.csv file.

1. Perform data cleaning procedures such that your final dataset is usable in the following questions. (2 points)

```
In [16]: import pandas as pd
import csv
import matplotlib.pyplot as plt

path = r'C:\Users\Shubhi Ambast\Desktop\DS_processedClevelandData.csv'
file = pd.read_csv(path)
```

```
In [17]: file.shape
```

```
Out[17]: (303, 14)
```

```
In [18]: file.isna().any() #To check NA or NaN value in the data
```

```
Out[18]: id           False
Age             False
Sex             False
cp              False
trestbps        False
chol            False
fbs             False
restecg         False
thalach         False
exang           False
slope           False
ca              False
thal            False
num             False
dtype: bool
```

```
In [19]: columns = ['ca', 'thal']
'''There are certain values in the above two features which does not contain any val
for col in columns:
    file[col] = file[col].str.replace('?', '0.0')
```

```
col_median = file[col].median()
file[col]=file[col].replace(to_replace = 0.0, value = col_median)
print(file)
```

	id	Age	Sex	cp	trestbps	chol	fb	restecg	thalach	exang	slope	\
0	1	63	1	1	145	233	1	2	150	0	3	
1	2	67	1	4	160	286	0	2	108	1	2	
2	3	67	1	4	120	229	0	2	129	1	2	
3	4	37	1	3	130	250	0	0	187	0	3	
4	5	41	0	2	130	204	0	2	172	0	1	
...
298	299	45	1	1	110	264	0	0	132	0	2	
299	300	68	1	4	144	193	1	0	141	0	2	
300	301	57	1	4	130	131	0	0	115	1	2	
301	302	57	0	2	130	236	0	2	174	0	2	
302	303	38	1	3	138	175	0	0	173	0	1	

	ca	thal	num
0	0	6	0
1	3	3	1
2	2	7	1
3	0	3	0
4	0	3	0
...
298	0	7	1
299	2	7	1
300	1	7	1
301	1	3	1
302	0.0	3	0

[303 rows x 14 columns]

1. For each type of diagnosis of heart disease, find the following for the resting blood pressure: (2 points)

- Mean
- Median
- Standard deviation

There are two types of diagnosis of heart disease:

- value 0 < 50% diameter narrowing
- value 1 > 50% diameter narrowing

```
In [20]: value = [0,1]
for i in value:
    num = file.loc[file['num'] == i]
    print('num value:',i)

    mean = num['trestbps'].mean()
    print('Mean:{:.2f}'.format(mean))

    median = num['trestbps'].median()
    print('Median:{:.2f}'.format(median))

    std = num['trestbps'].std()
    print('Std:{:.3f}'.format(std))
    print('\n')
```

```
num value: 0
Mean:129.25
Median:130.00
```

Std:16.205

```

num value: 1
Mean:134.57
Median:130.00
Std:18.769

```

1. Use Spearman's and Kendall correlation to quantify the correlation between age and the following.

- a. Resting blood pressure
- b. Serum cholesterol level
- c. Maximum heart rate achieved

Also, which variable(s) are most correlated with age? Illustrate with heatmaps. (3 points)

```

In [21]: age = file.iloc[:,1]
         bp = file.iloc[:,4]
         chol = file.iloc[:,5]
         heart_rate = file.iloc[:,8]

         corr = age.corr(bp,method = 'spearman')
         print('Correlation between age and resting blood pressure:',round(corr,3))
         corr = age.corr(chol,method = 'spearman')
         print('Correlation between age and Serum cholesterol level:',round(corr,3))
         corr = age.corr(heart_rate,method = 'spearman')
         print('Correlation between age and Maximum heart rate achieved:',round(corr,3))
         corr = age.corr(bp,method = 'kendall')
         print('Correlation between age and resting blood pressure:',round(corr,3))
         corr = age.corr(chol,method = 'kendall')
         print('Correlation between age and Serum cholesterol level:',round(corr,3))
         corr = age.corr(heart_rate,method = 'kendall')
         print('Correlation between age and Maximum heart rate achieved:',round(corr,3))

```

```

Correlation between age and resting blood pressure: 0.292
Correlation between age and Serum cholesterol level: 0.191
Correlation between age and Maximum heart rate achieved: -0.392
Correlation between age and resting blood pressure: 0.206
Correlation between age and Serum cholesterol level: 0.132
Correlation between age and Maximum heart rate achieved: -0.276

```

```

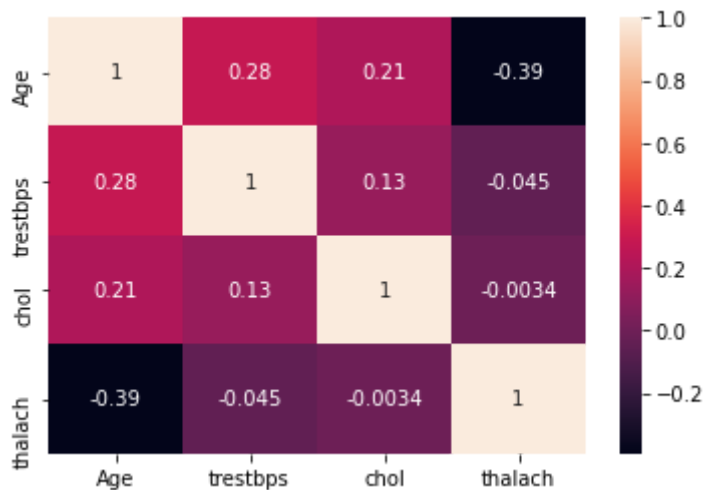
In [22]: import seaborn as sns
         data = file.drop(['id','Sex','cp','fbs','restecg','exang','slope','ca','thal','num']
         data.head(10)
         Var_Corr = data.corr()
         sns.heatmap(Var_Corr, xticklabels=Var_Corr.columns, yticklabels=Var_Corr.columns, an

```

```

Out[22]: <matplotlib.axes._subplots.AxesSubplot at 0x22f3a8a09a0>

```



It can be observed that resting blood pressure and serum cholesterol level are most correlated with age.


1. From your understanding, which of the features can be labeled as discrete random variables and which features as continuous random variables? (1 point)

Features which can be labelled as:

- discrete random variables : id, age, sex, ca
- continuous random variables : cp, trestbps, chol, fbs, restecg, thalach, exang, slope, thal, num

1. Describe the distribution for the values of the "thalach" feature? Illustrate with a plot. (1 point)

```
In [23]: labels = file['thalach']
labels.hist(bins=50, figsize=(10,5))
plt.show()
```




Thalach feature in the given dataset describes maximum heart rate achieved. It can be observed from the above graph that the frequency of heart rate lies in the range of 140-180 in the given dataset.

1. Plot the frequency of "Sex" variable in the dataset and describe what you observe in the

plot. Similarly plot and describe the 'ca' feature for the male participants. (2 points)


```
In [24]: sex = file['Sex']
sex.hist(bins=50, figsize=(10,5))
plt.show()
```



For the 'sex' plot, value = 0 for females and value = 1 for males.

The dataset has almost double frequency of males over females.

```
In [25]: male= file.loc[file['Sex'] == 1]
ca_male = male['ca']
ca_male.hist(bins=50, figsize=(10,5))
plt.show()
```




In the dataset, 'ca' tells about the number of major vessels labelled from 0-3. In males, the most common vessel is with can value 0

1. Detect outlier patients for features "trestbps" and "chol". Illustrate with plots. (2 points)

When plotting a boxplot, the outlier is the data point which is located outside the whiskers of box plot.


```
In [26]: import seaborn as sns
sns.boxplot(x=file['trestbps'])
```

Out[26]: <matplotlib.axes._subplots.AxesSubplot at 0x22f3aac93d0>



```
In [27]: import seaborn as sns
sns.boxplot(x=file['chol'])
```

Out[27]: <matplotlib.axes._subplots.AxesSubplot at 0x22f3a7e63a0>



Exercise 2 - Probability (total: 4 points)

1. Suppose a discrete random variable, MMSE (Mini mental state examination), cognitive test measured for Alzheimer's disease (AD) has the following probability mass function:

x	5	8	14	22	24	28	29	30
pr(X=x)	0.05	0.27	0.16	0.17	0.03	0.12	0.07	0.13

Find the probability that MMSE:

- at least 22 (1 point)
- at least 14 and at most 28 (1 point)

```
In [ ]: pmf_dict = {'5' : 0.05, '8':0.27, '14':0.16, '22':0.17, '24':0.03, '28':0.12, '29':0.07, '30':0.05, '32':0.05}

sum = 0
for i in pmf_dict.keys():
    if int(i) >= 22:
        sum += pmf_dict[i]
print("Probability that MMSE is atleast 22: {}".format(sum))

sum = 0
for i in pmf_dict.keys():
    if 14 <= int(i) <= 28 :
        sum += pmf_dict[i]
print("Probability that MMSE is atleast 14 and at most 28: {}".format(sum))
```

Probability that MMSE is atleast 22: 0.52

Probability that MMSE is atleast 14 and at most 28: 0.48

1. A company produced antibody testing kits for COVID-19. The false positive rate of the test is known to be 3%. What is the probability to find at least 2 false positive results within 35 tested patients? (2 points)

```
In [ ]: from scipy.stats import binom

""" Probability of finding atleast 2 positive = 1 - probability of finding 1 positive """

n = 35
p = 0.03
k = 1
probability = round(1 - binom.pmf(k,n,p),2)

print("Probability of finding atleast 2 false positive results is {}".format(probability))
```

Probability of finding atleast 2 false positive results is 0.63

Exercise 3 - Hypothesis Testing (total: 8 points)

Using the processed dataset from question 1 answer the following questions.

1. Are all the criteria for carrying out a t-test to identify a significant difference in the age of patients who have heart disease and those who don't, met? (3 points)
 - ☐ If the criteria is met, carry out a t-test using Python.
 - ☐ And if not, point out the unmet conditions for the variables, and mention a possible solution in-order to combat this issue.

```
In [3]: #Importing all Libraries necessary

import csv
import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns
%matplotlib inline
import scipy
import scipy.stats as stats
import math
from scipy.stats import f_oneway
```

```
from scipy.stats import ttest_ind
import warnings
warnings.filterwarnings('ignore')
```

In [4]: *#Opening and Reading CSV File*

```
PCD = pd.read_csv("DS_processedClevelandData.csv")
PCD.head(11)
```

Out[4]:

	id	Age	Sex	cp	trestbps	chol	fbs	restecg	thalach	exang	slope	ca	thal	num
0	1	63	1	1	145	233	1	2	150	0	3	0.0	6.0	0
1	2	67	1	4	160	286	0	2	108	1	2	3.0	3.0	1
2	3	67	1	4	120	229	0	2	129	1	2	2.0	7.0	1
3	4	37	1	3	130	250	0	0	187	0	3	0.0	3.0	0
4	5	41	0	2	130	204	0	2	172	0	1	0.0	3.0	0
5	6	56	1	2	120	236	0	0	178	0	1	0.0	3.0	0
6	7	62	0	4	140	268	0	2	160	0	3	2.0	3.0	1
7	8	57	0	4	120	354	0	0	163	1	1	0.0	3.0	0
8	9	63	1	4	130	254	0	2	147	0	2	1.0	7.0	1
9	10	53	1	4	140	203	1	2	155	1	3	0.0	7.0	1
10	11	57	1	4	140	192	0	0	148	0	2	0.0	6.0	0

In [5]: PCD.columns

Out[5]: Index(['id', 'Age', 'Sex', 'cp', 'trestbps', 'chol', 'fbs', 'restecg',
'thalach', 'exang', 'slope', 'ca', 'thal', 'num'],
dtype='object')

In [6]: print(PCD.shape)
print("-----", "\n")
print('Length of the data {}'.format(len(PCD)))
print("-----", "\n")
print(PCD.info())

(303, 14)

Length of the data 303

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   id           303 non-null    int64
1   Age          303 non-null    int64
2   Sex          303 non-null    int64
3   cp           303 non-null    int64
4   trestbps     303 non-null    int64
5   chol         303 non-null    int64
6   fbs         303 non-null    int64
7   restecg      303 non-null    int64
8   thalach      303 non-null    int64
9   exang        303 non-null    int64
10  slope        303 non-null    int64
11  ca           303 non-null    object
```



```

12 thal      303 non-null    object
13 num      303 non-null    int64
dtypes: int64(12), object(2)
memory usage: 33.3+ KB
None

```

```
In [7]: print (PCD.describe())
```

	id	Age	Sex	cp	trestbps	chol \
count	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000
mean	152.000000	54.438944	0.679868	3.158416	131.689769	246.693069
std	87.612784	9.038662	0.467299	0.960126	17.599748	51.776918
min	1.000000	29.000000	0.000000	1.000000	94.000000	126.000000
25%	76.500000	48.000000	0.000000	3.000000	120.000000	211.000000
50%	152.000000	56.000000	1.000000	3.000000	130.000000	241.000000
75%	227.500000	61.000000	1.000000	4.000000	140.000000	275.000000
max	303.000000	77.000000	1.000000	4.000000	200.000000	564.000000

	fbs	resetecg	thalach	exang	slope	num
count	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000
mean	0.148515	0.990099	149.607261	0.326733	1.600660	0.458746
std	0.356198	0.994971	22.875003	0.469794	0.616226	0.499120
min	0.000000	0.000000	71.000000	0.000000	1.000000	0.000000
25%	0.000000	0.000000	133.500000	0.000000	1.000000	0.000000
50%	0.000000	1.000000	153.000000	0.000000	2.000000	0.000000
75%	0.000000	2.000000	166.000000	1.000000	2.000000	1.000000
max	1.000000	2.000000	202.000000	1.000000	3.000000	1.000000

```
In [8]: PCD.Age.mean()
```

```
Out[8]: 54.43894389438944
```

Criteria:

1. Variables: We have One Numerical variable ("Age") and One Categorical Variable ("Num") to compare - So, we can use t-test for comparison of significant difference - we use two sample t-test as one is continuous variable and the other is qualitative variable.
2. Check Distribution of means of ages who have heart disease and who dont have
3. Calculating the mean, stdev.
4. Defining Hypothesis for t-test

Null Hypothesis: Average age of the person having a heart disease is less than the person who doesn't have heart disease.


Alternate Hypothesis: Average age of the person having a heart disease is more than the person who doesn't have heart disease.

```
In [9]: sns.set_style("ticks")

age_heart_disease = PCD.groupby('num')['Age']

fig, (axis1,axis2,axis3) = plt.subplots(1,3,figsize=(20,5))
ax = sns.distplot(PCD['Age'],ax=axis1)
ax.set(xlabel='Age')
ax = sns.distplot(age_heart_disease.get_group(0),ax=axis2)
ax.set(xlabel='Age With Heart Disease')
ax = sns.distplot(age_heart_disease.get_group(1),ax=axis3)
ax.set(xlabel='Age Without Heart Disease')
```

```
Out[9]: [Text(0.5, 0, 'Age Without Heart Disease')]
```



We can see the Mean distribution is different for both in above plots. So, we can perform t-test for Comparison.

```
In [10]: g_mean = PCD.groupby('Sex')['Age'].mean()
g_std = PCD.groupby('Sex')['Age'].std()


DataFrame = pd.DataFrame({'Group':[0,1], 'Sample_Age_Mean':g_mean.values, 'Sample_Age_Std':g_std.values, 'Sample_Size':[len(PCD.Age[PCD['Sex'] == 0]),len(PCD.Age[PCD['Sex'] == 1])])
DataFrame
```

```
Out[10]:
```

	Group	Sample_Age_Mean	Sample_Age_Std	Sample_Size
0	0	55.721649	9.370546	97
1	1	53.834951	8.836838	206

```
In [11]: sns.distplot(age_heart_disease.get_group(0),label='Presence_of_Heartdisease')
sns.distplot(age_heart_disease.get_group(1),label='Absence_of_Heart_disease')
plt.legend()
```

```
Out[11]: <matplotlib.legend.Legend at 0x7fd0ee4d1f70>
```



This above plot shows that that distribution of age of persons who doesnot have heart disease is lower than the distrubution of age who has heart disease.

```
In [12]: #Running T-Test from scipy library
```

```
In [13]: stats.ttest_ind(PCD.Age[PCD['num'] == 0],PCD.Age[PCD['num'] == 1],equal_var=False)
```

```
Out[13]: Ttest_indResult(statistic=-4.0303470663221095, pvalue=7.061439075547293e-05)
```

As already known, p value is 0.05, and for the current test, p value is less than 0.05 and so, null hypothesis is rejected and defined alternate hypothesis is true.

In []:

1. Identify if women are significantly more likely to get heart disease than men? (2 points)

In [14]:

```
PCD2 = PCD.copy()
```

In [15]:

```
def change_sex(Sex):
    if Sex == 0:
        return 'female'
    else:
        return 'male'
PCD2['Sex'] = PCD2['Sex'].apply(change_sex)
```


In [16]:

```
def presence(num):
    if num == 0:
        return "No Heart disease"
    else:
        return "Presence of Heart Disease"
PCD2['num'] = PCD2['num'].apply(presence)
```

In [17]:

```
sns.countplot(data=PCD2, x="Sex", hue="num")
plt.title('Gender v/s Heart Disease\n')
```

Out[17]: Text(0.5, 1.0, 'Gender v/s Heart Disease\n')



The above graph shows that the Women are not more likely to get Heart disease when compared to Men.

In []:

1. Inform yourself about χ^2 -test. And using χ^2 -test, identify if there is a significant association between exercise induced angina (exang), and the slope of the peak exercise ST segment (slope)? (3 points)

what is x square test

The aim of the test is to conclude whether there is an accociation between two variables or not we need to consider following things

The Contingency Table : A Contingency table (also called crosstab) is used to summarise the relationship between categorical variables.

Null hypothesis: We start by defining the null hypothesis (H_0) which states that there is no relation between the variables. An alternate hypothesis would state that there is a significant relation between the two.

```
In [ ]: # contingency table
contingency = pd.crosstab(PCD['exang'], PCD['slope'])
contingency
```

```
Out[ ]:  slope    1    2    3
      exang
      0  116   76   12
      1   26   64    9
```

```
In [ ]: from scipy.stats import chi2_contingency
        #It gives test statistics, the p-value, degrees of freedom and expected table
        c, p, dof, expected = chi2_contingency(contingency)
```

```
In [ ]: # interpret p-value
alpha = 0.05
print("p value is " + str(p))
#we need to compare the obtained p-value with alpha value
if p <= alpha:
    print('Dependent (reject H0)')
else:
    print('Independent (H0 holds true)')
```

```
p value is 3.490090811425977e-06
Dependent (reject H0)
```

As we reject H_0 , So, the variable exang have a significant association with variable slope.